


# NEURAL NETWORK FOR MACHINE LEARNING 2

한재근 과장 | [jahan@nvidia.com](mailto:jahan@nvidia.com)

유현곤 부장 | [hryu@nvidia.com](mailto:hryu@nvidia.com) / 양한별 과장 | [hanbyuly@nvidia.com](mailto:hanbyuly@nvidia.com)



# AGENDA

Training Neural Network

Stochastic Gradient Descent

Learning Rate Control

Data Preprocessing


Weight Initialization

Data Augmentation

Regularization & Dropout

Hyper parameter Exploration

# Training of Neural Networks

- 
1. Divide data on 2 sets: training set and validation set
  2. Define network and training procedure
    - network architecture
    - loss function
    - preprocessing and data augmentation
    - training algorithm and parameters (batch size, initialization,...)
  3. Training:  
for ( $t = 0$ ;  $t < T$ ;  $t++$ )  
    train  $\text{Net}(W)$  on training set
  4. test  $\text{Net}(W)$  on validation set  
If (not state-of-the art) goto 2

# Training of Neural Networks is difficult

A lot of hyper-parameters to choose:

- Loss function
- Network architecture: # of layers, and # of channels/layer
- Weight initialization
- SGD algorithm and learning rate policy
- Data preprocessing: scaling, augmentation ...

# Optimization Algorithms for CNN training

# Batch Gradient Descent

We want to minimize loss over training set with N samples  $(x_n, y_n)$ :

$$L(w) = \frac{1}{N} \sum_{n=1}^N E(f(x_n, w), y_n)$$

## Batch optimization:

1. accumulate gradients over all samples in training set

$$\frac{\partial E}{\partial y_{l-1}} = \frac{\partial E}{\partial y_l} \times \frac{\partial y_l(w, y_{l-1})}{\partial y_{l-1}} ; \quad \frac{\partial E}{\partial w_l} = \frac{\partial E}{\partial y_l} \times \frac{\partial y_l(w, y_{l-1})}{\partial w_l}$$

2. update W:

$$W(t+1) = W(t) - \lambda * \frac{1}{N} \sum_{n=1}^N \frac{\partial E}{\partial w} ((x_n, w), y_n)$$

## Issue:

Imagenet has  $10^6$  images  $\rightarrow$  gradient computation for whole set is expensive

# Stochastic Gradient Descent

## Stochastic Gradient Descent (on-line learning):

1. Randomly choose sample  $(x_k, y_k)$ :
2.  $W(t + 1) = W(t) - \lambda * \frac{\partial E}{\partial w}((x_k, w), y_k)$

## Stochastic Gradient Descent with mini-batches:


1. divide the dataset into small mini-batches, choose samples from different classes
2. compute the gradient using a single m-batch, make an update
3. move to the next mini-batch ...

**Don't forget to shuffle / shift data between epochs!**


# Stochastic Gradient Descent

$N = \text{batch}$ : gradient computation is heavy, step is small

$N = 1$  (on-line training): gradient is very noisy, zig-zags around “true” gradient



**batch**



**mini-batch**

Mini-batch training follows the curve of gradient:  
*the expected value of the weight change for on-line training is continuously pointing in the direction of the gradient at the current point in weight space.*



### Batch gradient descent

$$J_{train}(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Repeat {

$$\rightarrow \theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

$$\frac{\partial}{\partial \theta_j} J_{train}(\theta)$$

(for every  $j = 0, \dots, n$ )

}

### Stochastic gradient descent

$$\text{cost}(\theta, (x^{(i)}, y^{(i)})) = \frac{1}{2} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

$$J_{train}(\theta) = \frac{1}{m} \sum_{i=1}^m \text{cost}(\theta, (x^{(i)}, y^{(i)}))$$

1. Randomly shuffle dataset.

2. Repeat {

for  $i=1, \dots, m$  {

$$\theta_j := \theta_j - \alpha (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_j^{(i)}$$

(for  $j=0, \dots, n$ )

}

$(x^{(i)}, y^{(i)})$ .

$$\frac{\partial}{\partial \theta_j} \text{cost}(\theta, (x^{(i)}, y^{(i)}))$$

# Mini-batch Gradient Descent

## SGD with mini-batch:

1. faster than batch
2. more robust to redundant data
3. Behaves better in local minima/ saddle.



“Use stochastic gradient descent when training time is the bottleneck”  
*Leon Bottou, Stochastic Gradient Descent Tricks*

# Mini-batch size and Learning Rate

## How to change learning rate when we change the batch size?

**“Classical ML” rule:** on-line / mini-batch training with large learning rate is much more stable than batch training with the same learning rate.

**Convolutional NN** (for problems with large number of classes)

*Alex Krizhevsky (“One weird trick on parallelizing CNN”):*

1. **Theory:** multiply the learning rate by  $k^{1/2}$  when increase the batch size by  $K$  to keep the variance in the gradient expectation constant.
2. **Practice:** to multiply the learning rate by  $k$  when multiplying the batch size by  $k$ .

**WARNING:**

This rule does not work when mini-batch size become too large !

# Example: CIFAR0-10

airplane



automobile



bird



cat




deer



dog



frog



horse



ship




truck



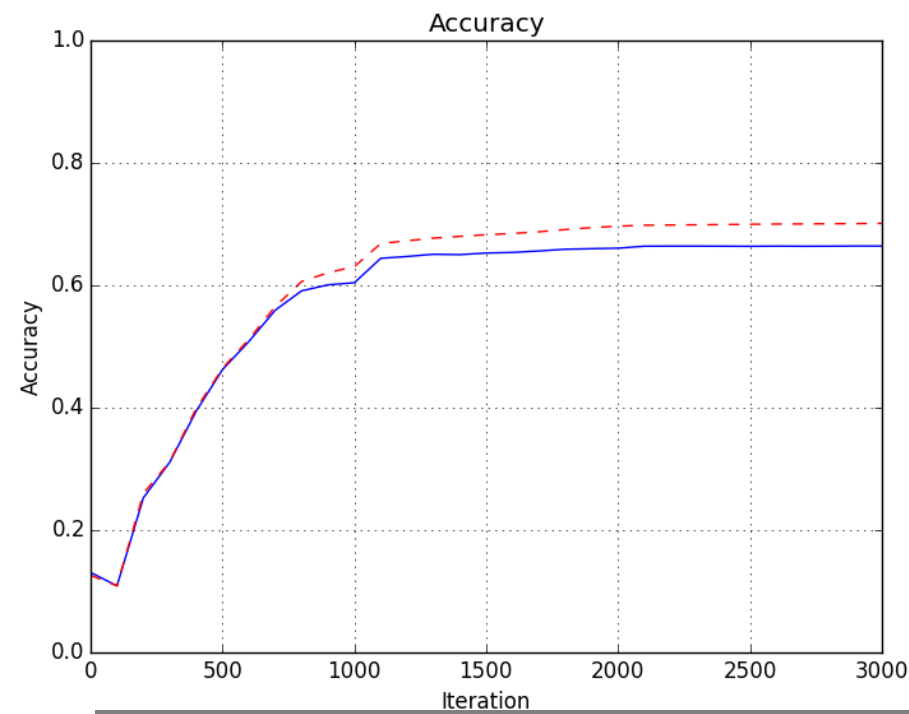
# Learning Rate & its control

# Example: CIFAR-10 training

Training and testing accuracy




Batch =100, initial lr =0.001




Batch =1000, initial lr =0.01

# Example: CIFAR-10

Training and testing accuracy



Batch =100, initial lr =0.001



Batch =1000, initial lr =0.001

# Learning Rate Adaptation

$$W(t + 1) = W(t) - \lambda(t) * \frac{\partial E}{\partial w}$$

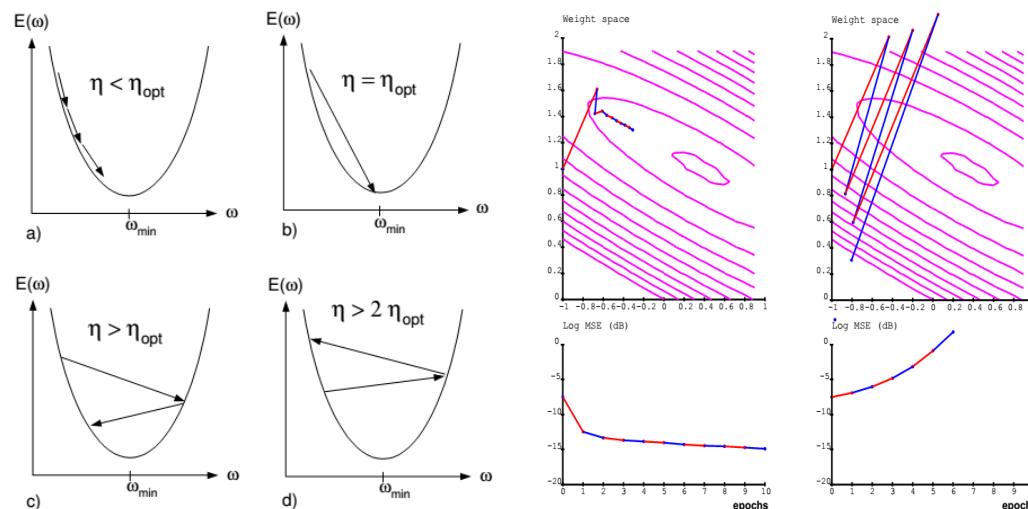
Classical learning rate annealing:

$$\sum_{t=1}^{\infty} \frac{1}{\lambda^2(t)} < \infty \quad \text{and} \quad \sum_{t=1}^{\infty} \frac{1}{\lambda(t)} = \infty ,$$

e.g.  $\lambda(t) = \frac{c}{t}$ ;


Other learning rate policies:

1. Manual:  $\lambda = \text{const}$
2. exp:  $\lambda_n = \lambda_0 * \gamma^n$
3. step :  $\lambda_n = \lambda_0 * \gamma^{\lceil \frac{n}{\text{step}} \rceil}$
4. inverse:  $\lambda_n = \lambda_0 * (1 + \gamma * n)^{-c}$





# Example: CIFAR-10 training




Initial lr = 0.001, decrease lr by 10x for 10,000 and 20,000 iterations

# Learning Rate

No silver bullet for learning rate

Choose proper size with explanation

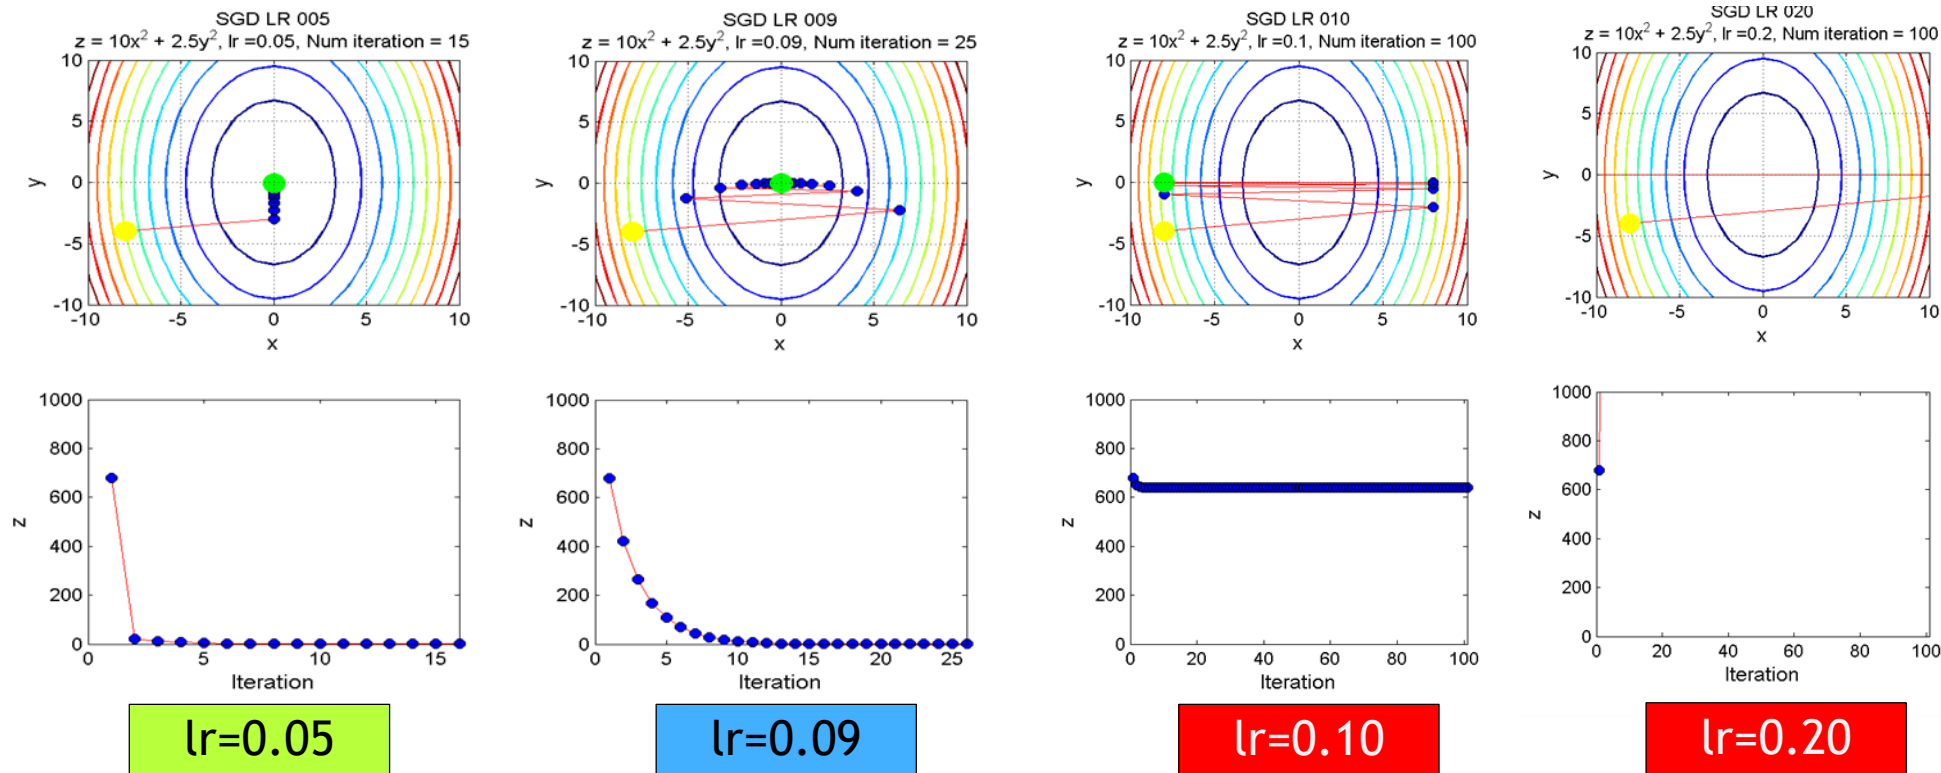


# Why accuracy improves when we decrease lr?

## Motivational example

Let's take a quadratic function  $z = 10 * x^2 + 2.5 * y^2$  and use SGD to find minimum.

Basic SGD converges if learning rate  $< \frac{1}{10}$ .



# Adaptive learning rate

## Adapt

1. learning rate per weight or per layer
  - Use only the sign of the gradient
  - Divide the learning rate for a weight by a running average of the magnitudes of recent gradients for that weight (Adagrad, AdaDelta, RMSPROP,..)
2. learning rate and momentum
  - Natural gradient
  - ADAM (*Knigmn and BA*)

# SGD with weight decay

$$W(t + 1) = W(t) - \lambda * \left( \frac{\partial E}{\partial w} + \theta * W(t) \right)$$

Weight decay works to keep weights size under control (“regularization”).


This is equivalent to adding penalty on weights to loss function:

$$E'(W) = E(W) + \frac{\theta}{2} * W^2$$

# SGD with momentum

$$W(t + 1) = W(t) + \Delta W(t + 1)$$

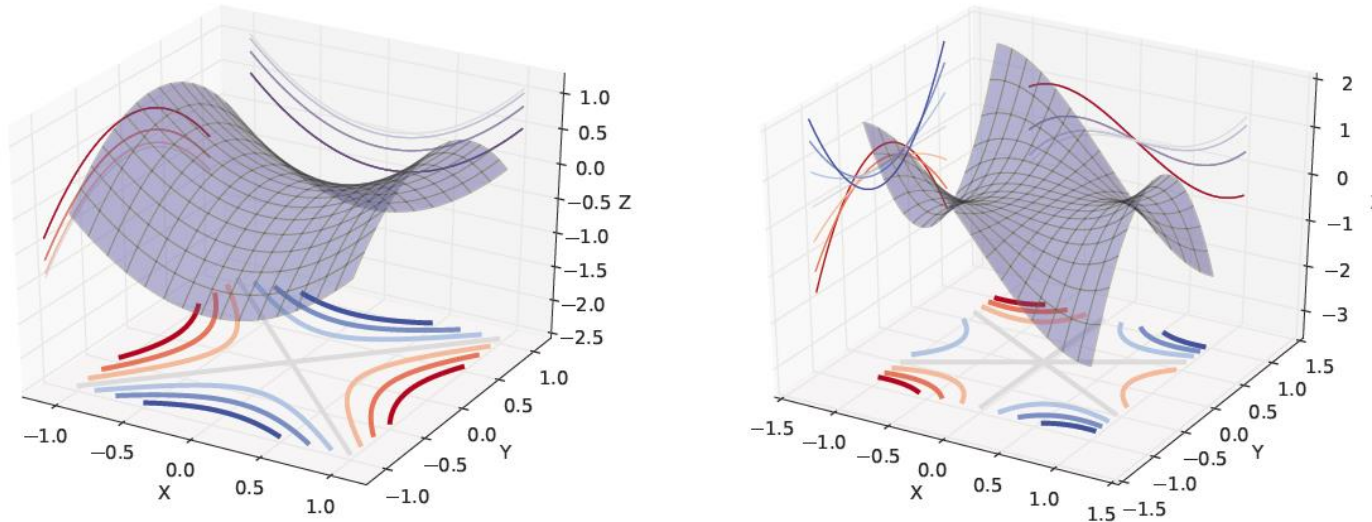
$$\Delta W(t + 1) = \beta * \Delta W(t) - \lambda * \frac{\partial E}{\partial w}$$



Momentum works as weighted average of gradients .

$$\Delta W(t+1) = -\lambda * \left( \sum_{k=1}^{t+1} \beta^{t+1-k} * \frac{\partial E}{\partial w}(k) \right)$$

# Optimization near Saddle Points



“The problem with convnets cost functions is not local min, but local saddle points. How SGD methods behave near saddle point?”

R. Pascanu, “On the saddle point problem for non-convex optimization”, <http://arxiv.org/abs/1405.4604>

Dauphin, “Identifying and attacking the saddle point problem in high-dimensional non-convex optimization”, <http://arxiv.org/pdf/1406.2572v1.pdf>

# Advanced Optimization Algorithms

A lot of interesting optimization algorithms to speed-up training and get better results:

- Nesterov Accelerated Gradient
- Adagrad
- Adadelta
- RMSPROP/RPROP
- Variance-based SGD
- Averaged SGD
- ADAM



# AdaGrad, AdaDelta, and ADAM

**Adagrad:** adapt learning rate for each weight

$$\Delta W_{ij}(t+1) = -\frac{\gamma}{\sqrt{\sum_1^{t+1} (\frac{\partial E}{\partial w_{ij}}(\tau))^2}} * \frac{\partial E}{\partial w_{ij}}(t+1)$$



**AdaDelta:** accumulate the denominator over last k gradients (sliding window):

$$\alpha(t+1) = \sum_{t-k+1}^{t+1} (\frac{\partial E}{\partial w_{ij}}(\tau))^2$$
$$\Delta W_{ij}(t+1) = -\frac{\gamma}{\sqrt{\alpha(t+1)}} * \frac{\partial E}{\partial w_{ij}}(t+1)$$

This requires to keep k gradients. Instead we can use simpler formula:

$$\beta(t+1) = \rho * \beta(t) + (1 - \rho) * (\frac{\partial E}{\partial w_{ij}}(t+1))^2$$
$$\Delta W_{ij}(t+1) = -\frac{\gamma}{\sqrt{\beta(t+1) + \epsilon}} * \frac{\partial E}{\partial w_{ij}}(t+1)$$


# Performance near Saddle Points



by Alec Radford [alecradford](https://www.github.com/alecradford) <http://imgur.com/a/Hqolp>


# Data Preprocessing

# Preprocessing the data



# Dimension Reduction

In practice, you may also see **PCA** and **Whitening** of the data




# TLDR: In practice for Images: center only

e.g. consider CIFAR-10 example with [32,32,3] images

- Subtract the mean image (e.g. AlexNet)  
(mean image = [32,32,3] array)
- Subtract per-channel mean (e.g. VGGNet)  
(mean along each channel = 3 numbers)

Not common to normalize  
variance, to do PCA or  
whitening

# Case: DeepFace Architecture




Yaniv Taigman, etc (Facebook) . DeepFace: Closing the Gap to Human-Level Performance in Face Verification, CVPR 2014

# Weight Initialization



# Initialization for Neural Network



```
W = 0.01* np.random.randn(D,H)
```







# Gradient banishment

```
W = 0.01* np.random.randn(D,H)
```

```
W = np.random.randn(fan_in, fan_out) * 1.0 # layer initialization
```


input layer had mean 0.000927 and std 0.998388  
 hidden layer 1 had mean -0.000117 and std 0.213081  
 hidden layer 2 had mean -0.000001 and std 0.047551  
 hidden layer 3 had mean -0.000002 and std 0.010630  
 hidden layer 4 had mean 0.000001 and std 0.002378  
 hidden layer 5 had mean 0.000002 and std 0.000532  
 hidden layer 6 had mean -0.000000 and std 0.000119  
 hidden layer 7 had mean 0.000000 and std 0.000026  
 hidden layer 8 had mean -0.000000 and std 0.000006  
 hidden layer 9 had mean 0.000000 and std 0.000001  
 hidden layer 10 had mean -0.000000 and std 0.000000

input layer had mean 0.001800 and std 1.001311  
 hidden layer 1 had mean -0.000430 and std 0.981879  
 hidden layer 2 had mean -0.000849 and std 0.981649  
 hidden layer 3 had mean 0.000566 and std 0.981601  
 hidden layer 4 had mean 0.000483 and std 0.981755  
 hidden layer 5 had mean -0.000682 and std 0.981614  
 hidden layer 6 had mean -0.000401 and std 0.981560  
 hidden layer 7 had mean -0.000237 and std 0.981520  
 hidden layer 8 had mean -0.000448 and std 0.981913  
 hidden layer 9 had mean -0.000899 and std 0.981728  
 hidden layer 10 had mean 0.000584 and std 0.981736




# Understanding the difficulty of training convolutional networks


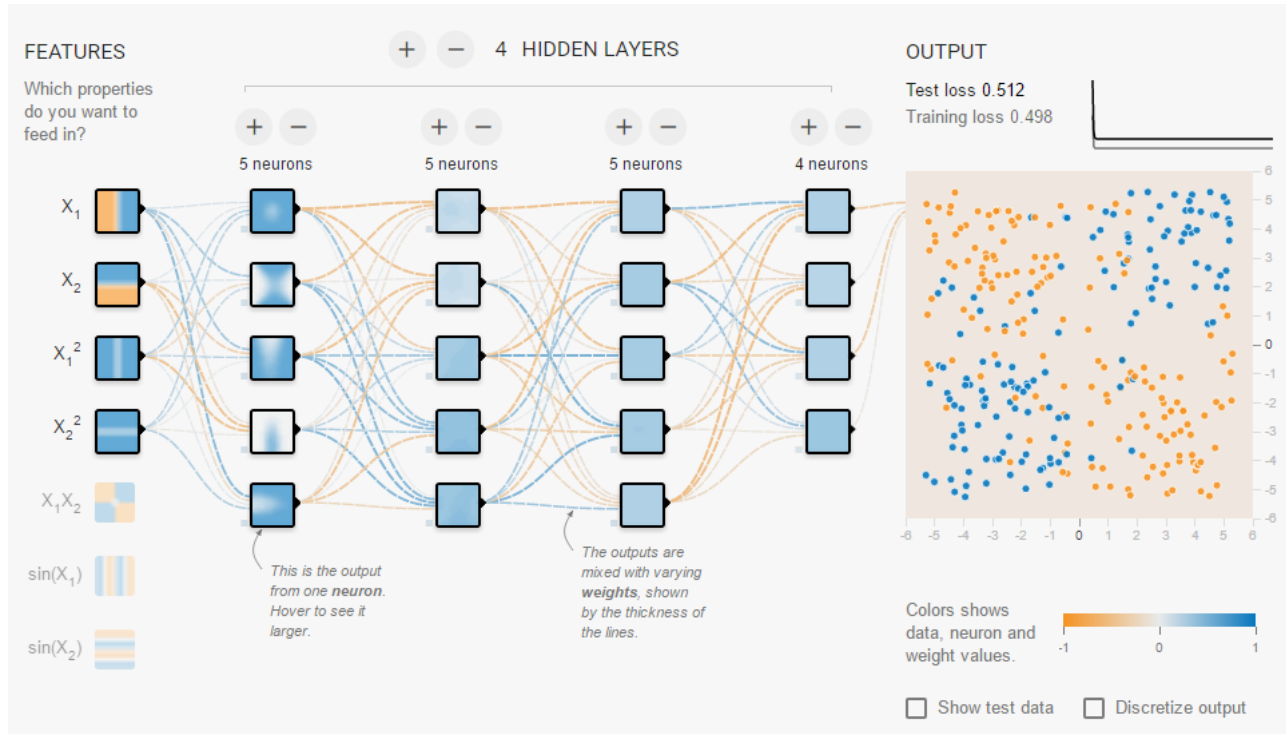
Example: MLP with 4 fully connected layers, with sigmoid non-linear function.  
Measure mean and standard deviation of the activation (output of the sigmoid) for 4 hidden layers



The top hidden layer quickly saturates at 0. It drives the gradient  $dE/dy$  multiplied by 0 and therefore causing the gradient descent to stall, slowing down all learning. Near epoch 100 it slowly de-saturates releasing other layers



# Why?






**Sigmoid**


All neurons will be all zero during back-propagation

# Xavier initialization

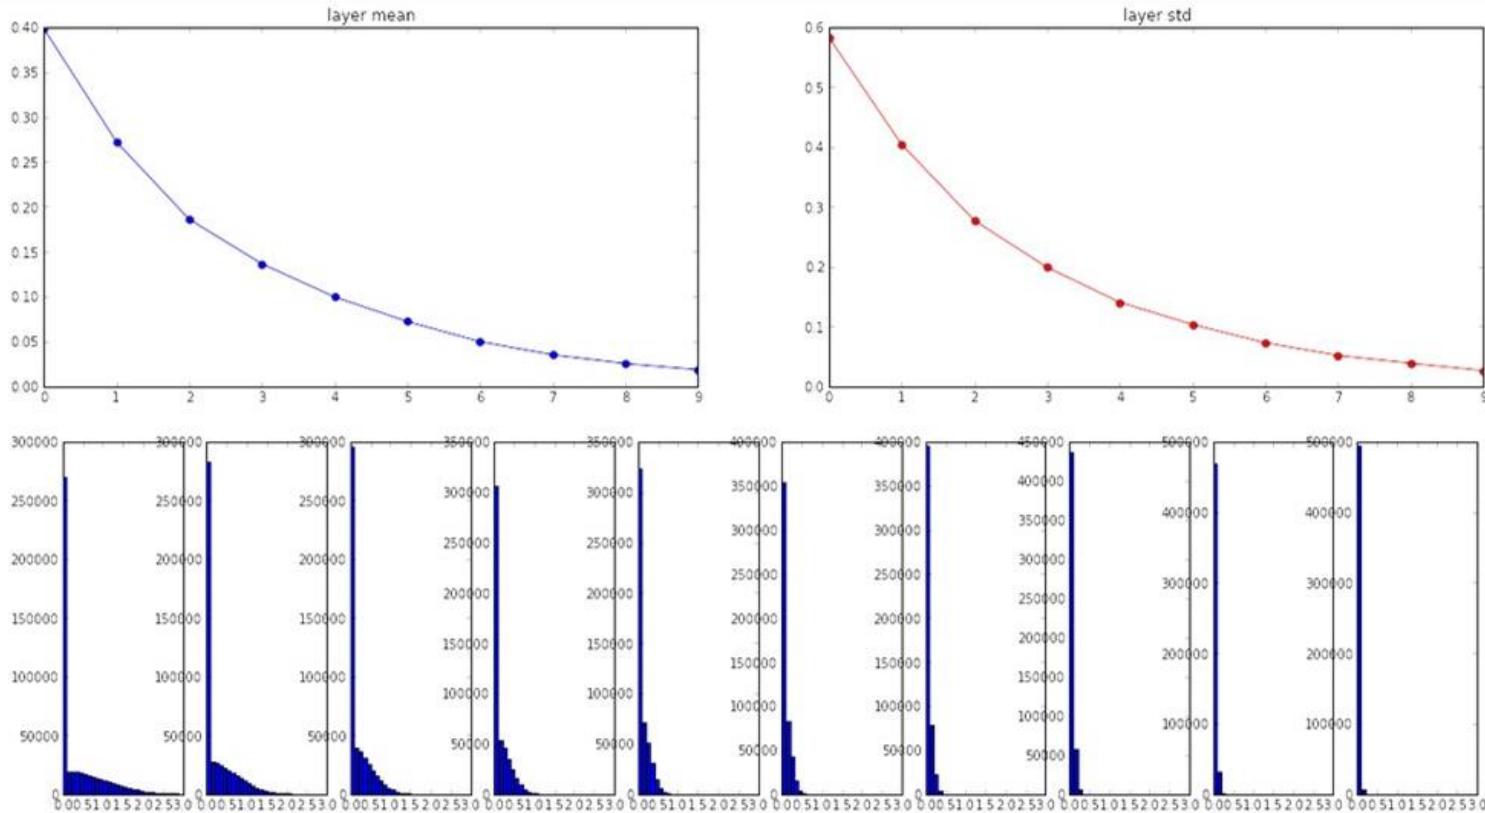
```
W = np.random.randn(fan_in, fan_out) / np.sqrt(fan_in) # layer initialization
```



**Reasonable initialization.**  
(Mathematical derivation  
assumes linear activations)

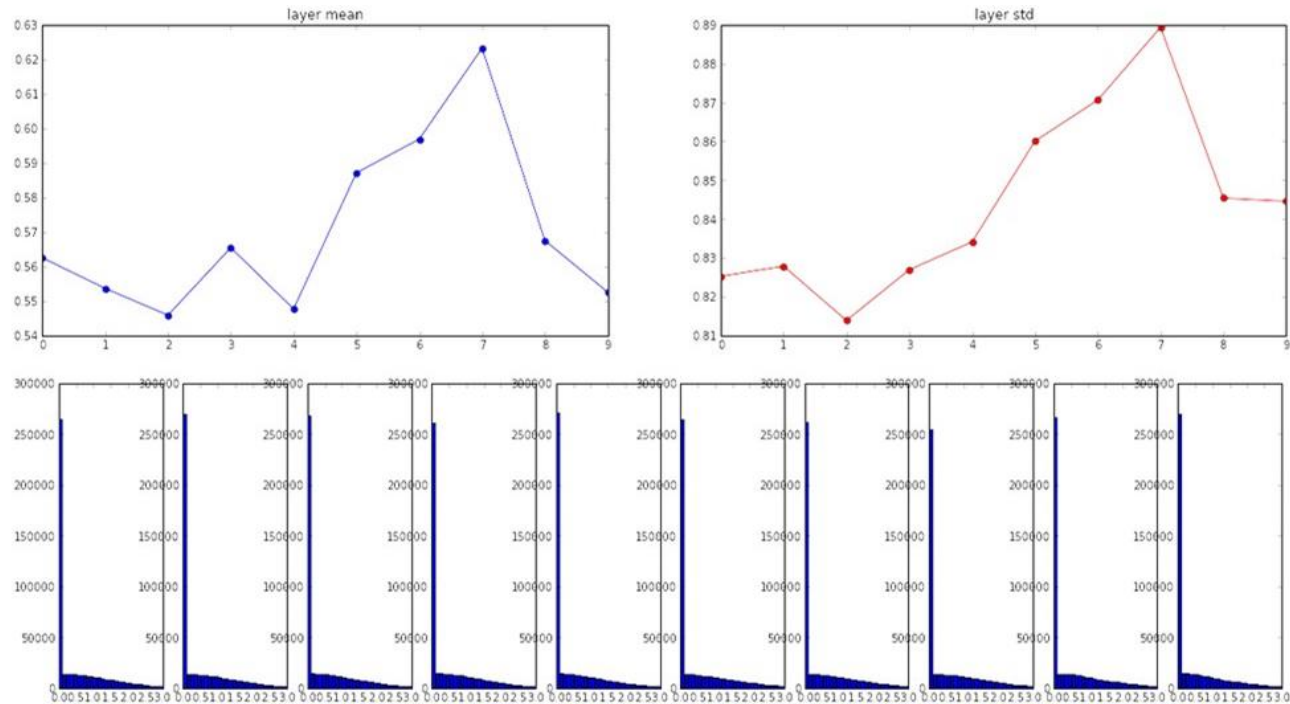


# Xavier Initialization with ReLU



# He initialization

```
W = np.random.randn(fan_in, fan_out) / np.sqrt(fan_in/2) # layer initialization
```



# Regularization | Data Augmentation



# Neural Network Regularization

- ▶ Empirical
  - ▶ Dropout
  - ▶ DropConnect
  - ▶ Stochastic pooling
  - ▶ Artificial Data
- ▶ Explicit
  - ▶ Early Stopping
  - ▶ Limiting Number of Parameters
  - ▶ Weight Decay
    - ▶ L1/L2 regularization
  - ▶ Max norm constraints

# Data Augmentation

The common method to “enlarge” training set:

- image translations
- re-scale (both up and down) before crop
- horizontal and vertical reflections ( “flip”)
- elastic deformation with random interpolations ((bilinear, area, nearest neighbor and cubic, with equal probability) (*Simard, 2003*))
- photometric distortion and altering the intensities of the RGB channels in training images (A.G. Howard., 2013)

# Data Augmentation (1)

## 1. Horizontal Flips



## 2. Random crops/scales

**Training:** sample random crops / scales


ResNet:

1. Pick random L in range [256, 480]
2. Resize training image, short side = L
3. Sample random 224 x 224 patch

**Testing:** average a fixed set of crops

ResNet:

1. Resize image at 5 scales: {224, 256, 384, 480, 640}
2. For each size, use 10 224 x 224 crops: 4 corners + center, + flips



# Data augmentation

## 3. Color Jitter




## 4. Others

Random mix/combinations of :

- translation
- rotation
- stretching
- shearing,
- lens distortions, ... (go crazy)

# Data augmentation (Plankton competition)




<http://benanne.github.io/2015/03/17/plankton.html>


# Regularization Dropout

# Dropout Layer

randomly set some neurons to zero in the forward pass



(a) Standard Neural Net



(b) After applying dropout.

# Code example

```
p = 0.5 # probability of keeping a unit active. higher = less dropout

def train_step(X):
    """ X contains the data """


    # forward pass for example 3-layer neural network
    H1 = np.maximum(0, np.dot(W1, X) + b1)
    U1 = np.random.rand(*H1.shape) < p # first dropout mask
    H1 *= U1 # drop!

    H2 = np.maximum(0, np.dot(W2, H1) + b2)
    U2 = np.random.rand(*H2.shape) < p # second dropout mask
    H2 *= U2 # drop!

    out = np.dot(W3, H2) + b3

    # backward pass: compute gradients... (not shown)
    # perform parameter update... (not shown)
```



Example forward pass with a 3-layer network using dropout






# Why this is good?

It helps higher accuracy




(a) Without dropout



(b) Dropout with  $p = 0.5$ .

# Using for Training / Inference



```
def train_step(X):
```

```
    # forward pass for example 3-layer neural network
```

```
    H1 = np.maximum(0, np.dot(W1, X) + b1)
```

```
    U1 = (np.random.rand(*H1.shape) < p) / p # first dropout mask. Notice /p!
```


```
    H1 *= U1 # drop!
```

```
    H2 = np.maximum(0, np.dot(W2, H1) + b2)
```

```
    U2 = (np.random.rand(*H2.shape) < p) / p # second dropout mask. Notice /p!
```

```
    H2 *= U2 # drop!
```

```
    out = np.dot(W3, H2) + b3
```



```
def predict(X):
```

```
    # ensembled forward pass
```

```
    H1 = np.maximum(0, np.dot(W1, X) + b1) # no scaling necessary
```

```
    H2 = np.maximum(0, np.dot(W2, H1) + b2)
```


```
    out = np.dot(W3, H2) + b3
```

# Hyperparameter Optimizations

# Parameters affecting training

- Learning rate
- Regularizations
- Filter size
- Model depth
- Stride
- Dropout
- Model architecture
- ...

# Random Search vs. Grid Search



*Random Search for Hyper-Parameter Optimization*  
Bergstra and Bengio, 2012

# Search optimal hyper-parameters

## example

```
max_count = 100
for count in xrange(max_count):
    reg = 10**uniform(-5, 5)
    lr = 10**uniform(-3, -6)
```

adjust range

```
max_count = 100
for count in xrange(max_count):
    reg = 10**uniform(-4, 0)
    lr = 10**uniform(-3, -4)
```

```
val_acc: 0.527000, lr: 5.340517e-04, reg: 4.097824e-01, (0 / 100)
val_acc: 0.492000, lr: 2.279484e-04, reg: 9.991345e-04, (1 / 100)
val_acc: 0.512000, lr: 8.680827e-04, reg: 1.349727e-02, (2 / 100)
val_acc: 0.461000, lr: 1.028377e-04, reg: 1.220193e-02, (3 / 100)
val_acc: 0.460000, lr: 1.113730e-04, reg: 5.244309e-02, (4 / 100)
val_acc: 0.498000, lr: 9.477776e-04, reg: 2.001293e-03, (5 / 100)
val_acc: 0.469000, lr: 1.484369e-04, reg: 4.328313e-01, (6 / 100)
val_acc: 0.522000, lr: 5.586261e-04, reg: 2.312685e-04, (7 / 100)
val_acc: 0.530000, lr: 5.808183e-04, reg: 8.259964e-02, (8 / 100)
val_acc: 0.489000, lr: 1.979168e-04, reg: 1.010889e-04, (9 / 100)
val_acc: 0.490000, lr: 2.036031e-04, reg: 2.406271e-03, (10 / 100)
val_acc: 0.475000, lr: 2.021162e-04, reg: 2.287807e-01, (11 / 100)
val_acc: 0.460000, lr: 1.135527e-04, reg: 3.905040e-02, (12 / 100)
val_acc: 0.515000, lr: 6.947668e-04, reg: 1.562808e-02, (13 / 100)
val_acc: 0.531000, lr: 9.471549e-04, reg: 1.433895e-03, (14 / 100)
val_acc: 0.509000, lr: 3.140888e-04, reg: 2.857518e-01, (15 / 100)
val_acc: 0.514000, lr: 6.438349e-04, reg: 3.033781e-01, (16 / 100)
val_acc: 0.502000, lr: 3.921784e-04, reg: 2.707126e-04, (17 / 100)
val_acc: 0.509000, lr: 9.752279e-04, reg: 2.850865e-03, (18 / 100)
val_acc: 0.500000, lr: 2.412048e-04, reg: 4.997821e-04, (19 / 100)
val_acc: 0.466000, lr: 1.319314e-04, reg: 1.189915e-02, (20 / 100)
val_acc: 0.516000, lr: 8.039527e-04, reg: 1.528291e-02, (21 / 100)
```

**53%** - relatively good for a 2-layer neural net with 50 hidden neurons.

But this best cross-validation result is worrying. Why?

# Monitor and visualize the loss curve

